



速成教育
SPEED UP EDUCATION



CSCB09

Software Tools and Systems Programming

导师： VC

UTSC Week 02 Class | 2022/1/17

S P E E D U P E D U C A T I O N

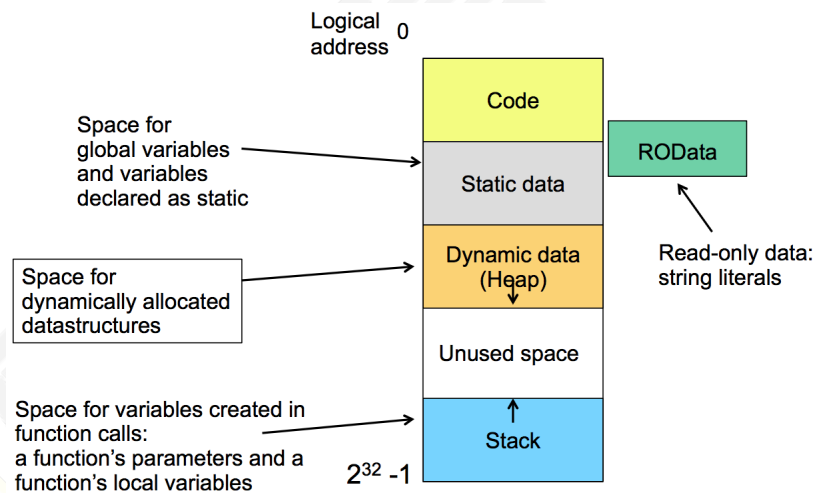
Disclaimer

This study handout is provided by Speed Up Education and its affiliated tutors. The study handout seeks to support your study process and should be used as a complement, not substitute to the university course material, lecture notes, problem, sets, past tests and other available resources.

The study handout is distributed for free to the students participating in Speed Up Education classes and is not for sale or other commercial uses whatsoever. We kindly ask you to refrain from copying or selling in part or in whole any information provided in the handout.

Thanks for choosing Speed Up Education. We wish you the best of luck in your exam.

Memory



Type	Size	Value Range
char	1 byte	-128 to 127 or 0 to 255
unsigned char	1 byte	0 to 255
signed char	1 byte	-128 to 127
int	2 or 4 bytes	-32,768 to 32,767 or -2,147,483,648 to 2,147,483,647
unsigned int	2 or 4 bytes	0 to 65,535 or 0 to 4,294,967,295
short	2 bytes	-32,768 to 32,767
unsigned short	2 bytes	0 to 65,535
long	8 bytes or (4bytes for 32 bit OS)	-9223372036854775808 to 9223372036854775807
unsigned long	8 bytes	0 to 18446744073709551615
pointer	8 bytes	...

File I/O in C

File functions work with `FILE *` all the time:

`FILE` = Type representing stream state. Actual definition varies by compiler and OS, but most likely a **struct**, so we pass around **pointers**.

(Poor naming because it isn't about files, it's about contents and streams.

'STREAM' would be better. I should also say "stream functions" instead of "file functions".)

Header to `#include: <stdio.h>`

Open

```
FILE *fopen(const char *filename, const char *mode)
```

Mode	meaning	if file exists	if not
"r"	read	cool	error
"w"	write	truncate	create
"a"	append	cool	create
"r+"	read+write	cool	error
"w+"	write+read	truncate	create
"a+"	append+read	cool	create

"read" and "write" start at the beginning of the file.

"append" start at the end of the file.

"truncate" = erase original content.

Returns NULL if error. (This one you should always check.)

Close

```
int fclose(FILE *stream)
```

Returns 0 if success, EOF if error.

Why need to close:

1. Limit on how many streams are open per process.
2. Writing may be delayed (buffered) until closing. (More on buffering later.)
3. Windows: No two processes can open the same file.

Formatted I/O

```
int fprintf(FILE *stream, const char *format, ...)
int fscanf(FILE *stream, const char *format, ...)
```

In fact:

```
printf(format, args) = fprintf(stdout, format, args)
scanf(format, args) = fscanf(stdin, format, args)
```

`stdin` and `stdout` are pre-opened streams for standard-in and standard-out.

There is one more: `stderr`, so you can output error messages and leave 'stdout' for normal data/output.

(No need to manually close 'stdin', 'stdout', 'stderr'.)

* Usually just use 'perror' to print error message to stderr:

```
void perror(const char *prefix)
```

Character I/O

One single character:

```
int putchar(int c) /* stdout */
int putc(int c, FILE *stream)
int getchar(void) /* stdin */
int getc(FILE *stream)
```

Returns the character written/read if success, `EOF` if error or end of stream.

Important:

`EOF` fits in 'int' but not 'char'. Correct usage of 'getchar' and 'getc' involves:

1. Store return value in 'int' variable, not 'char'.
2. Test for `EOF`.
3. If not `EOF`, safe to down-convert to 'char'.

String I/O

```
int fputs(const char *string, FILE *stream)
```

Returns EOF if error. Also difference from 'puts': 'fputs' does not add newline at the end, 'puts' does.

```
char *fgets(char *dest, int n, FILE *stream)
```

Reads at most $n - 1$ characters or until (and including) newline.

Returns NULL if error or end of stream, dest if success.

Exercise 1: Why does it need you to provide n ?

Exercise 2: Why $n - 1$ characters read, not n ?

Arbitrary Data I/O

```
size_t fread(void *dest, size_t s, size_t n, FILE *stream)
```

```
size_t fwrite(const void *data, size_t s, size_t n, FILE *stream)
```

Read/Write n items, each item s bytes (ask `sizeof`); in-memory raw bytes used.

Returns how many items read/written.

Some use cases:

- A whole array.

- Record (struct) (single or array of).

- Raw bytes (array of 'unsigned char' usually).

Cross-platform watchout:

The same raw bytes can mean different values on different platforms. Check compatibility!

Seeking

```
int fseek(FILE *stream, long int i, int origin)
```

Set the file pointer to the location $\text{origin} + i$

origin	go to i bytes from:
SEEK_SET	beginning
SEEK_END	end
SEEK_CUR	current position

Returns non-zero on error.

Error vs End-of-Stream Disambiguation

`getc` returns EOF if end-of-stream or error.
`fgets` returns NULL.
`fread` returns $< n$.
etc.

How to tell end-of-stream from error:

```
int feof(FILE *stream) Returns true if end-of-stream status.
```

Important:

Does **not** predict for next read; only **remembers** for previous read. You must read before asking.

```
int ferror(FILE *stream) Returns true if error status.
```

Buffering

C file I/O delays writing:

accumulates data in buffer until large chunk, then requests kernel to write that chunk.

And hastens reading:

requests kernel to read a large chunk into buffer, then serves your puny 'getc' etc. from buffer.

Why:

Huge overhead per kernel request (system call, "syscall") regardless of data size.

1 syscall for 1000 bytes beats 1000 syscalls for 1 byte each.

But can be disabled/reconfigured if your application needs.